


Mindfully Yours — End-to-End AI Architecture (Phase 1: English only)

Per-user isolation · DPDP-compliant, India · Phase 1 launches in English — Hindi added in Phase 2 (see roadmap panel) · two-tier caching for cost control · every AI task is retrieved, triaged, guarded, reviewed and logged



Dashed arrows: Storage → Runtime (reads scoped data) · Prompt engineering → Runtime (loads governed prompt into assembly) · Feedback loop ↔ Prompt engineering (rollback cycle) · Feedback loop ↔ Observability (shared signals)

Cost control: two-tier cache (FAQ/exact-match → semantic) checked before any retrieval or LLM call — safety checks never skipped, even on a cache hit · common TTS phrases cached as audio

Phase 1 scope: English only, end to end — conversation, safety interlock vocabulary, and voice all in English. The same Sarvam model stack extends to Hindi in Phase 2 without an architecture change.

Also in scope: streaming at every stage (target ~1.5–2.5s) · graceful degradation avatar → voice → text on provider outage · AI service falls back to interlock-only safe path if the model API is unavailable

Latency masking: on a slower turn, brief natural filler phrases ("let me think about that...", "okay...") play while the real response is still being generated — so a moment of processing time never feels like a stall to the user

Mindfully Yours — AI-Led Mental & Emotional-Wellness Triage Platform · Prepared by OpenXcell Technolabs Pvt. Ltd.

Sections ⑩ and ⑪ are engineering recommendations for the technology review, not yet client-confirmed scope. The Hindi query-handling panel presents two valid options for that same review.