

Mindfully Yours — Data Usage Map

19 — Data Usage Map: Which File Feeds Which Part of the System

Purpose: a single, permanent reference for implementation. Every real data source this project has received — the original 6 Phase I clinical spreadsheets and the September 2026 "convo data" drop — mapped to exactly one of three usage buckets, with the reasoning for that placement and the concrete mechanism that consumes it. Read this before writing any ingestion or prompt-assembly code — it answers "which table does this file become, and what touches it at runtime" in one place, so nobody has to re-derive it file-by-file.

This document assumes familiarity with [04_CONVERSATION_PIPELINE.md](#) (the runtime pipeline this feeds), [05_KNOWLEDGE_BASE_DEEP_DIVE.md](#) (Phase I sheet detail), and [18_NEW_CONVO_DATA_CATALOG.md](#) (new-drop detail) — this file is the cross-cutting index over both, not a replacement for either.

The one rule that decides everything below

Ask one question about any piece of content: **is it something a person might say, a real thing to hand back to them, or a number/rule the code just checks?**

If it is...	It becomes...	It is never...
Real, varied human phrasing (something a user might say, or the AI's own tone/style)	Bucket 1 or 2 — embedded for matching, or shown once as a style example	A hardcoded rule
A real exercise, script, or piece of guidance worth giving back to a user	Bucket 2 — searched for and retrieved	Baked into the prompt permanently
A count, a threshold, a fixed decision sequence	Bucket 3 — a code-level rule, checked directly	Embedded or searched at all

Bucket 1 — Teach the AI how to talk (style only; loaded once, never searched live)

File (exact name)	Why it belongs here	How it's actually used
Convo_Samples.pdf	Real example conversations showing tone/pacing across common scenarios — the point is how it's said, not looking one up mid-chat	A curated subset becomes fixed few-shot examples inside the system prompt (see 17_PERSONA_AND_CONVERSATION_STYLE.md); the rest is a held-out evaluation set, never live-retrieved

File (exact name)	Why it belongs here	How it's actually used
Career and Academic Conversation - Samples.docx , Conversation Combinations - Tanisha.docx , Conversation Combinations- Anam.docx , Conversations - Anxiety .docx , Conversations - Depression.docx , ED/Eating Disorder - Third Person Reported.docx , Miscellaneous Conversations1.docx , 18-9_Offline Woork Week 6.docx , and the Tanisha/ Ankita/Shreya weekly conversation-script files under MYPL KB 7th Sept/	Same reasoning as Convo_Samples.pdf — real dialogue showing tone, especially valuable because several of these reach genuine crisis moments the original sample file never did	Curate a fresh few-shot set from these (esp. the crisis-handling examples) and fold into the same persona prompt; the bulk becomes an expanded evaluation set
Clinical Markers Reference Table (1).pdf	Only 10 items, always relevant, needs to be visible on literally every turn — chunking/ searching it would be solving a problem it doesn't have	Lives permanently in the triage system prompt, verbatim (see app/safety/clinical_markers.py in the backend) — never embedded, never a retrieval target
Career Indecision - Canonical Template (deduplicated).md	A style/pattern example (career-indecision opening, including the deliberate "I won't give concrete career advice" scope boundary) — deduplicated down from Career Conversation Sets 1.docx 's 500 copy-pasted instances to the 1 real template it actually was	Loaded once as a few-shot example alongside the other career-scenario samples; its [CAREER_NAME] / [STREAM] placeholders are filled from the Bucket 2a lookup CSV below, not generated by the LLM

Bucket 2 — Search for it, then either match against it or hand it back

This bucket has **two different jobs** that must not be confused: **(a) matching** — figuring out what the user means — and **(b) recommending** — finding real content to give them. Both are "search," but the output is different.

2a. Matching (detecting what's going on — nothing is shown to the user directly)

File (exact name)	Why it belongs here	How it's actually used
Kb_Phase_I_A1.csv (Signal Legend) — the phrasings column specifically	285 signals, each with 5-10 real example phrasings — exactly the varied, real-world language a similarity search needs	Embedded; the user's message is compared against these embeddings to detect which clinical signal(s) are present (Step A of 04_CONVERSATION_PIPELINES.md)
Kb_Phase_I_A2.csv (Symptom Map) — the how_described column	Same reasoning as A1 — real phrasing worth matching against, for the 21 rows that genuinely have it (see 05_KNOWLEDGE_BASE_DEEP_DIVE.md §3 for why the other 133 rows don't)	Embedded and matched the same way as A1
Kb_Phase_I_9.csv (Signal Sufficiency) — the pattern_name field only	The only narrow case worth embedding on this sheet — a rough early guess at which of ~185 patterns applies, before the exact signal checklist takes over	Embedded; used for a first coarse match, then the exact minimum/must-have signal check (Bucket 3) takes over
Kb_Phase_I_7.csv (Red Flags) — the trigger_phrases column	Real examples, deliberately written to include disguised/indirect language — a genuine "does this match, by meaning" problem	Embedded with a deliberately more permissive match threshold than normal (see 08_SAFETY_INTERLOCK_AND_TRIAGE.md §4) — but see the special case below for what happens after a match
PHASE II_Emotion Associated user expression.docx , List of Emotions - Tanisha.docx , and the Pathway... doc's emotion picklist — once merged into one canonical list (see §Blockers below)	Real first-person "I feel..." phrasing, meant to teach/pattern-match emotion detection	Embedded and matched against a user's message to tag which emotion is present
List of Scenarios - Tanisha.docx , Anxiety Situations.docx	Master scenario/topic taxonomies — matching "which life situation is this" is a real semantic-similarity problem, not a lookup	Embedded; used to route a conversation toward the right topic-specific question bank or conversation pattern
Layer I Questions and Statments.docx , Layer II Question Bank.docx , Layer III.docx , PHASE II Layer I and IV Questions.docx , Layer 2 to 3 Transitions.docx	These are a menu to select the single best-fitting phrase from, given the current conversational depth — a retrieval problem, not free generation	Embedded per layer; the pipeline picks the closest-fitting question/statement for "what layer are we in right now" and hands it to the LLM as phrasing material (never the LLM inventing the question from scratch)
Career Field & Stream Lookup.csv (deduplicated from Career Conversation Sets 1.docx — see below)	500 real (career name, field, academic stream) rows — exactly the kind of exact, structured lookup that should be matched by name, not generated	Matched when a user names a specific career; returns the field category and stream needed to fill the Bucket 1 template's placeholders

2b. Recommending (real content handed back to the user)

File (exact name)	Why it belongs here	How it's actually used
<code>dbt-skills-workbook.docx</code> / <code>dbt-skills-workbook-simple.pdf</code>	A finished, ready-to-use DBT skills tool — literally the "self-care tool" content 04_CONVERSATION_PIPELINE.md Step F was designed to search for and previously had nothing to find	Chunked (one skill = one chunk), embedded, and genuinely semantically searched at Tier 1 — the one part of the whole system where full RAG search is used exactly as designed
<code>brain-dump-worksheets.pdf</code>	Same reasoning — finished, ready-to-hand-out content	Same mechanism as the DBT workbook
<code>SELF-HELP CONVERSATIONS.docx</code> (the toolkit half — Behavioral Activation, Time Management, Mindfulness/Grounding, Journaling, Tracking Worksheets sections)	Broadest single self-help asset received — real, usable exercises	Same mechanism; each named tool/worksheet becomes its own searchable chunk
Phase III/IV content in general, once the client's remaining gaps are filled (see 15_OPEN_QUESTIONS_AND_BLOCKERS.md #24)	This is the category Phase III (clinical guidance/psycho-education) and Phase IV (self-help tools) were always meant to populate	Same mechanism — this bucket grows as more finished content arrives

Bucket 3 — Becomes literal code logic (never text, never embedded, never searched)

File (exact name)	Why it belongs here	How it's actually used
<code>Kb_Phase_I_C11.csv</code> (Routing Rules)	A fixed, precise signal-combination → tier mapping — searching for it would be slower and less reliable than reading it off a table	Loaded as a direct lookup table, evaluated strictly in <code>priority</code> column order (1→2→3), never <code>rule_id</code> order (see 05_KNOWLEDGE_BASE_DEEP_DIVE.md §4) — <code>app/safety/routing.py</code> in the backend

File (exact name)	Why it belongs here	How it's actually used
Kb_Phase_I_9.csv (Signal Sufficiency) — every field except <code>pattern_name</code> (<code>minimum_signals</code> , <code>must_have_signals</code> , <code>confidence_rule</code> , <code>graceful_exit_trigger</code>)	Exact lookup criteria, checked once a candidate pattern is already known from Bucket 2a's coarse match	Read directly off the row matched in 2a — no second search
Kb_Phase_I_17.csv (Fallback & Safety Net) — the <code>trigger_condition</code> field	Describes the conversation's own shape (turn count, response length, confidence score) — not something a user says, so there's no phrase to search for	Evaluated as code against the running conversation state (Redis) every turn — app/safety/fallback.py
Pathway For AI In Different Situations - Tanisha.docx	This is, almost verbatim, a 15-step decision flow with explicit branches — the closest thing to pseudocode the client has produced	Becomes the actual control flow of app/pipeline/orchestrator.py — identify concern → categorize → safety check → assess emotion/severity → pick intervention → close. Its vocabulary (Severity buckets, the 22-item category list) still needs reconciling with the existing Tier/scenario systems first — see 15_OPEN_QUESTIONS_AND_BLOCKERS.md #23
Knowledge_Base_Directory	Internal build-planning document only	Never reaches the AI at all — not ingested in any form
Filler-phrase library (Mindfully_Yours_AI_Latency_Filler_Phrases.pdf)	Selection is a rule match (situation type + expected gap length → category), not a meaning-based search	A lookup table in app/pipeline/filler.py , picked by rule, never embedded

Special case: exact-fetch content (a hybrid of Bucket 2's matching and Bucket 3's "no thinking allowed")

Some content is found via a Bucket 2a match, but the response itself must be used **word-for-word**, never paraphrased by the LLM — this is its own category precisely because getting it wrong is the single biggest safety risk in the system:

File (exact name)	Field	Rule
Kb_Phase_I_7.csv (Red Flags)	bot_script	Matched via embedding (Bucket 2a), then fetched verbatim — the LLM is never called for this response. See 08_SAFETY_INTERLOCK_AND_TRIAGE.md §4 for the 4 flags (RF-028-031) that currently have no script at all and must hard-alert a human instead.
Kb_Phase_I_17.csv (Fallback & Safety Net)	bot_script	Same rule — fetched verbatim once a fallback scenario is confirmed by Bucket 3's code check.

Career Conversation Sets 1.docx — fixed, not excluded

Was nominally "500 conversations," actually 1 template repeated 500 times with only the career name and academic stream swapped. **Deduplicated, not discarded** — split into its two real parts, now correctly bucketed:

- Bucket 1: [Career Indecision - Canonical Template \(deduplicated\).md](#) — the one real template.
- Bucket 2a: [Career Field & Stream Lookup.csv](#) — the 500 genuine (career, field, stream) rows, usable to fill the template's placeholders or to match a user-stated career interest to its field/stream. Both files live under [documnets/The Real Useful Knowledge/](#).

What's explicitly NOT usable yet (don't wire these in)

File (exact name)	Status	Why
Anam/Conversation Combinations-sleep.docx , Anam/Conversations Combinations- anxiety 3rd person.docx , Anam/Miscellaneous Conversations.docx	Still password-locked	Didn't decrypt with the password used for the rest of the drop — need a second password from Anam
Paraphilic-disorder content in List of Scenarios - Tanisha.docx ("Fetishes" section) and the Week 2 conversation-script file	On hold	Needs an explicit clinical + legal governance decision before it goes anywhere near a live index — see 15_OPEN_QUESTIONS_AND_BLOCKERS.md #26
Conversations - Depression.docx , Conversations Combination (Bipolar and Psychosis).docx , Shreya/Trauma Related Conversations.docx , Conversation Combinations- Life Transitions.docx , Ankita/Conversations around Parenting.docx	Partially usable	Each has real, specific content gaps (missing scenarios, cuts off mid-file) — usable for what's actually there, but don't treat as complete; see 15_OPEN_QUESTIONS_AND_BLOCKERS.md #24

Blockers this data-usage map depends on being resolved first

Before Bucket 2a's emotion-matching and Bucket 3's Pathway-as-code work can be built for real (not just structurally stubbed), these need a decision — full detail in

[15_OPEN_QUESTIONS_AND_BLOCKERS.md](#):

- **#23** — merge the three unreconciled emotion lists (19/9/11 labels) into one canonical taxonomy, and reconcile the new Severity (Mild/Moderate/Severe/Crisis) scale and 22-item category list against the existing Tier 1/2/3 and ~9-category systems.
- **#25** — standardize tier-labeling, since different clinical-team contributors used it inconsistently.

Companion to [04_CONVERSATION_PIPELINE.md](#), [05_KNOWLEDGE_BASE_DEEP_DIVE.md](#), and [18_NEW_CONVO_DATA_CATALOG.md](#). Update this table the moment a new file arrives or a blocker above is resolved — this is meant to be the one place a developer checks before deciding how to treat a new piece of content, not a snapshot to go stale.