

# Mindfully Yours — Database Tables, Explained Simply

## 23. Database Tables, Explained Simply

One-page reference: every table in the database, what it stores, why it exists, and when it actually gets used. Written for a non-technical read — no code, just plain language.

There are **19 tables** in total. They fall into 6 natural groups, by job, not by when they were built.

### Quick-look summary

| Table                    | What it's for                               | Ever vector-searched?             |
|--------------------------|---------------------------------------------|-----------------------------------|
| kb_red_flags             | Danger phrases + pre-written safety replies | <b>Yes — working today</b>        |
| kb_content_chunks        | Self-help tools to recommend                | <b>Yes — working today</b>        |
| kb_signals               | Fine-grained clinical warning signs         | Designed for it, not wired up yet |
| kb_symptoms              | Body/emotion symptoms → real cause          | Designed for it, not wired up yet |
| kb_sufficiency_patterns  | "Do we know enough yet?" thresholds         | Designed for it, not wired up yet |
| kb_scenarios             | Real-life situation matching                | Designed for it, not wired up yet |
| kb_emotion_phrases       | Which emotion a message expresses           | Designed for it, not wired up yet |
| kb_layer_content         | How deep to go in conversation              | Designed for it, not wired up yet |
| kb_clarification_phrases | "I'm not sure, let me ask" phrases          | Designed for it, not wired up yet |
| kb_career_lookup         | Career name → field & stream                | Designed for it, not wired up yet |
| kb_tier_question_bank    | Real follow-up questions                    | Designed for it, not wired up yet |
| kb_anxiety_situations    | Anxiety topic checklist                     | No — plain lookup                 |
| kb_fallback_scenarios    | What to do when things go off-script        | No — plain lookup                 |
| kb_routing_rules         | "If X, always route to Y"                   | No — plain lookup                 |
| kb_guardrail_rules       | Parenting-concern urgency                   | No — plain lookup                 |
| kb_versions              | Tracks knowledge-base updates               | No — system bookkeeping           |
| prompt_config            | The AI's fixed instructions                 | No — system bookkeeping           |
| sessions                 | One row per conversation                    | No — system bookkeeping           |
| turn_audit_log           | One row per message, for audits             | No — system bookkeeping           |

"Designed for it, not wired up yet" means: the data is loaded and each row already has its numeric fingerprint saved — the table is ready. What's still missing is the search code on the other end that would actually query it during a live conversation. Only two tables have that search code written today.

## Group 1 — Safety (checked on every single message, no exceptions)

### kb\_red\_flags — 31 rows

- **Stores:** a danger-sign ID, ~5 example phrases of what someone in that danger might say, what it clinically means, what to do immediately, and the exact pre-written reply to send.
- **Why:** so the AI never has to improvise when someone's safety is at risk — it just matches and reads a pre-approved script, word for word.
- **Used:** every single message, instantly, before anything else in the pipeline runs.

### kb\_signals — 285 rows

- **Stores:** a clinical concept (e.g. "hopelessness"), how severe it typically is (1-3), whether it counts as a red flag, and example phrasings.
- **Why:** the AI's early-warning system — a much finer-grained read of "what's going on" than the 31 red flags alone.
- **Used:** meant to run on every message to build a running picture of the user's state. (Built, but the live query isn't wired up yet — see summary table.)

### kb\_symptoms — 154 rows

- **Stores:** physical/emotional symptoms (e.g. "chest tightness"), how people typically describe them, and which [kb\\_signals](#) rows they usually point back to.
- **Why:** users often describe a body feeling, not a clinical term — this bridges "my chest feels tight" to the real underlying cause.
- **Used:** alongside kb\_signals, same not-yet-wired status.

### kb\_fallback\_scenarios

- **Stores:** what to do when a conversation isn't going normally — user gives too little information, contradicts themselves, goes quiet, or belongs to a special group (e.g. a minor).
- **Why:** a script for "this isn't a normal exchange," so the AI doesn't freeze or guess.
- **Used:** whenever the conversation flow breaks the usual pattern.

## Group 2 — Understanding what the user means

### kb\_scenarios

- **Stores:** a category, a topic, a short scenario description, and an example phrase someone might actually say.
- **Why:** matches vague or indirect language to a known real-life situation (relationship trouble, career stress, etc.).
- **Used:** to classify what kind of situation the user is describing.

### kb\_emotion\_phrases

- **Stores:** an emotion name (Overwhelm, Grief, Anxiety, ...) and one real example sentence per row.
- **Why:** tells the AI what emotion sits behind a message, even when the wording is indirect.
- **Used:** reading the emotional tone of a message.

## kb\_layer\_content

- **Stores:** which "depth level" (Layer I-IV) a question or statement belongs to, and its text.
- **Why:** conversations go deeper in stages — this tracks how deep this exchange should go, separate from how clinically severe it is.
- **Used:** deciding how probing versus surface-level the next question should be.

## kb\_clarification\_phrases

- **Stores:** a category and a ready-made phrase to use when the AI genuinely isn't sure what the user means.
- **Why:** gives the AI a safe way to ask for clarification instead of quietly guessing.
- **Used:** any low-confidence moment, at any point in the conversation.

## kb\_anxiety\_situations

- **Stores:** a broad anxiety situation (e.g. "Public speaking") and its related sub-topics, as one block of text.
- **Why:** a coverage checklist, so the team can confirm anxiety topics are represented — not built for fine-grained matching.
- **Used:** as a reference list, not a live search target.

# Group 3 — Recommending something to do

## kb\_content\_chunks

- **Stores:** a piece of real help content (a DBT skill, a self-care tool), which tier it applies to, and its full text.
- **Why:** this is the one table whose content genuinely gets handed back to a struggling user as a suggestion.
- **Used:** Tier 1 conversations, whenever the AI offers a coping tool.

## kb\_sufficiency\_patterns

- **Stores:** a clinical pattern name, which tier (1-3) it maps to, and how much evidence counts as "enough" to make that call.
- **Why:** stops the AI from deciding how serious something is on a hunch — it's a checklist of what counts as sufficient evidence.
- **Used:** every message, to decide whether enough is known yet to assign a tier.

# Group 4 — Career-specific lookups

## kb\_career\_lookup

- **Stores:** a career name, its broader field, and its academic stream.
- **Why:** a precise reference table so career-planning answers are factual, never invented.
- **Used:** whenever a user names or asks about a specific career.

## kb\_tier\_question\_bank

- **Stores:** real follow-up questions (open-ended or multiple-choice), tied to a scenario or career topic.
- **Why:** gives the AI real, pre-written questions to ask next, instead of making its own up.

- **Used:** during scenario-specific conversations, to choose the next question.

## Group 5 — Pure rules (never searched — just looked up, like a checklist)

### kb\_routing\_rules

- **Stores:** a trigger condition and where to route the conversation, in a fixed priority order.
- **Why:** deterministic — "if X happens, always do Y," no AI judgment needed or wanted.
- **Used:** deciding where a conversation should go next.

### kb\_guardrail\_rules

- **Stores:** parenting/child-related concern categories and how urgently each one needs a human referral.
- **Why:** same idea as routing rules, scoped specifically to parenting topics.
- **Used:** checked directly in code whenever a parenting concern comes up.

## Group 6 — System bookkeeping (not conversation content at all)

### kb\_versions

- **Stores:** a label and a status ("building" / "live" / "rolled\_back") for each batch of knowledge-base data that gets loaded.
- **Why:** lets a knowledge-base update be tested before it goes live, and safely rolled back if something's wrong with it.
- **Used:** every time the knowledge base is updated.

### prompt\_config

- **Stores:** the AI's fixed persona and instruction text, versioned, with a clinical sign-off record.
- **Why:** the prompt is never editable by any user or admin on the fly — this is the one governed, auditable place it's allowed to live.
- **Used:** loaded once at the start of every conversation, into every prompt sent to the AI.

### sessions

- **Stores:** one row per conversation — its mode (voice/text/avatar), language, current tier, a running summary, and a mood tag.
- **Why:** remembers who the user is and what's happened so far, across an entire conversation (and into the next one).
- **Used:** from the start to the end of every conversation.

### turn\_audit\_log

- **Stores:** one row per message — what was retrieved, what matched, what tier was assigned, and how long each step took.
- **Why:** a legally required record, so anyone can later check exactly why the AI said what it said.
- **Used:** written silently after every message, never slowing the user down while it happens.