


Today's Build, Explained Simply

The one-sentence version

Until today, the AI could talk — but it never actually looked anything up. Now it does: it reads your message, checks it against the doctors' real data, decides how serious it is using their real rules, and only then lets the AI compose a reply — using a real doctor-written instruction, not a guess.

1. Before vs. after — the biggest change of all

The AI was never even told what you said. A bug meant every single reply was generated from a generic instruction sheet only, with your actual message missing entirely.



2. How a message gets handled now — 4 steps

Everything new fits into one simple recipe, run once per message:

1. ONE EMBED

Turn the message into one numeric "fingerprint" — **exactly once** per message.

The only slow, costly step (~0.3–0.6s, a network call) — everything below reuses this same fingerprint for free.



2. FAN OUT

Check that fingerprint against several doctor-written tables at once — cheap and fast, since it's just your own database, not the internet.

kb_red_flags	danger phrases (31 rows)
kb_signals	clinical warning signs (285 rows)
kb_symptoms	body/emotion descriptions (154 rows)
kb_sufficiency_patterns	closest matching clinical pattern (195 rows)



3. ACCUMULATE

Don't trust one message alone — add whatever matched to a running tally for the **whole conversation**, kept in memory across every message.

A single match is a clue, not a conclusion — this is what stops the system from jumping to a diagnosis after one ambiguous sentence.



4. GATE

Only ever hand the AI ONE of two safe things:

- **A real safety script**, word for word — if something dangerous matched
- **A doctor-written question to ask** — never a diagnosis label, and only once enough evidence has built up

The AI's only job past this point is phrasing it kindly — it never decides what's true.

3. A real example — same message, before and after

Message: "I recently moved to a totally new city on my own and I feel really isolated, can't seem to make any friends here."

Before today, the reply would have been a generic opening line — because the AI never actually saw this message at all, in any form.

After today, here's what really happened:

What was checked	What it found
Danger phrases	Nothing dangerous — safe to continue
Clinical signals recognized	8 real matches, including "Leaving Home" and "Disproportionate distress"
Rule looked up	A real doctor-written rule matched
Seriousness decided	Tier 2 — worth real attention, not an emergency

The actual reply the AI gave, grounded in that real match, asked about how long it's been going on, whether it's affecting sleep/eating/work, what usually helps, and whether anything else happened around the same time — a genuine, structured check-in, not an improvised guess. All of this took about **1.7 seconds**, comfortably inside the 1.5–2.5 second target.

4. Nine real bugs, found by actually testing it

Code review wouldn't have caught these — only running the system against real messages did.

#	What was wrong	Why it mattered
1	The user's message was never in the prompt	Every reply was generic, regardless of what was typed
2	A doctor's real routing notes were silently empty	A tricky punctuation mark meant the computer couldn't find the column
3	Two unrelated things looked like a match	Two different IDs happened to share the same number

#	What was wrong	Why it mattered
4	"Missing" data wasn't actually missing	It was sitting in a different table the whole time
5	A speed measurement was counting the same thing twice	Made a slow step look faster than it really was
6	The chat AI's connection was set up wrong	Real conversation replies were failing silently
7	The chat AI occasionally sent an empty reply chunk	Crashed the whole response
8	The safety record-keeping was quietly failing	Every conversation's audit trail was being lost
9	A "how serious is this" result wasn't being saved	Made it look like a check never ran, when it had

5. What's genuinely working now, verified live

- A dangerous message still gets the real safety response, instantly, no AI improvisation.
- A safe, happy message still gets a warm, natural reply — and now it's actually about what you said.
- A message describing real distress gets matched against real clinical data and a real doctor-written follow-up question — not a guess.
- All of this stays inside the speed budget.

6. What's honestly still not done

- Recommending an actual self-care exercise (Tier 1) has the code built, but hasn't been seen firing yet in testing — it needs a very specific combination to trigger.
- The AI's "style samples" (~20 example conversations, written to teach tone) still aren't loaded in.
- A chunk of the doctors' data (life-stage "layers," career/emotion scenario libraries) is still waiting on the client to confirm how it should combine with severity levels — not an engineering gap, a pending decision.

Nothing here has been committed to git yet — still holding for a go-ahead, per your instruction.