

Mindfully Yours — Data Usage Map

19 — Data Usage Map: Which File Feeds Which Part of the System

Purpose: a single, permanent reference for implementation. Every real data source this project has received — the original 6 Phase I clinical spreadsheets and the September 2026 "convo data" drop — mapped to exactly one of three usage buckets, with the reasoning for that placement and the concrete mechanism that consumes it. Read this before writing any ingestion or prompt-assembly code — it answers "which table does this file become, and what touches it at runtime" in one place, so nobody has to re-derive it file-by-file.

This document assumes familiarity with [04_CONVERSATION_PIPELINE.md](#) (the runtime pipeline this feeds), [05_KNOWLEDGE_BASE_DEEP_DIVE.md](#) (Phase I sheet detail), and [18_NEW_CONVO_DATA_CATALOG.md](#) (new-drop detail) — this file is the cross-cutting index over both, not a replacement for either.

The one rule that decides everything below

Ask one question about any piece of content: **is it something a person might say, a real thing to hand back to them, or a number/rule the code just checks?**

If it is...	It becomes...	It is never...
Real, varied human phrasing (something a user might say, or the AI's own tone/style)	Bucket 1 or 2 — embedded for matching, or shown once as a style example	A hardcoded rule
A real exercise, script, or piece of guidance worth giving back to a user	Bucket 2 — searched for and retrieved	Baked into the prompt permanently
A count, a threshold, a fixed decision sequence	Bucket 3 — a code-level rule, checked directly	Embedded or searched at all

Bucket 1 — Teach the AI how to talk (style only; loaded once, never searched live)

File (exact name)	Why it belongs here	How it's actually used
Convo_Samples.pdf	Real example conversations showing tone/pacing across common scenarios — the point is how it's said, not looking one up mid-chat	A curated subset becomes fixed few-shot examples inside the system prompt (see 17_PERSONA_AND_CONVERSATION_STYLE.md); the rest is a held-out evaluation set, never live-retrieved

File (exact name)	Why it belongs here	How it's actually used
Career and Academic Conversation - Samples.docx , Conversation Combinations - Tanisha.docx , Conversation Combinations- Anam.docx , Conversations - Anxiety .docx , Conversations - Depression.docx , ED/Eating Disorder - Third Person Reported.docx , Miscellaneous Conversations1.docx , 18-9 Offline Woork Week 6.docx , and the Tanisha/Ankita/Shreya weekly conversation-script files under MYPL KB 7th Sept/	Same reasoning as Convo_Samples.pdf — real dialogue showing tone, especially valuable because several of these reach genuine crisis moments the original sample file never did	Curate a fresh few-shot set from these (esp. the crisis-handling examples) and fold into the same persona prompt; the bulk becomes an expanded evaluation set
Clinical Markers Reference Table (1).pdf	Only 10 items, always relevant, needs to be visible on literally every turn — chunking/ searching it would be solving a problem it doesn't have	Lives permanently in the triage system prompt, verbatim (see app/safety/clinical_markers.py in the backend) — never embedded, never a retrieval target

Bucket 2 — Search for it, then either match against it or hand it back

This bucket has **two different jobs** that must not be confused: **(a) matching** — figuring out what the user means — and **(b) recommending** — finding real content to give them. Both are "search," but the output is different.

2a. Matching (detecting what's going on — nothing is shown to the user directly)

File (exact name)	Why it belongs here	How it's actually used
Kb_Phase_I_A1.csv (Signal Legend) — the phrasings column specifically	285 signals, each with 5-10 real example phrasings — exactly the varied, real-world language a similarity search needs	Embedded; the user's message is compared against these embeddings to detect which clinical signal(s) are present (Step A of 04_CONVERSATION_PIPELINES.md)
Kb_Phase_I_A2.csv (Symptom Map) — the how_described column	Same reasoning as A1 — real phrasing worth matching against, for the 21 rows that genuinely have it (see 05_KNOWLEDGE_BASE_DEEP_DIVE.md §3 for why the other 133 rows don't)	Embedded and matched the same way as A1

File (exact name)	Why it belongs here	How it's actually used
Kb_Phase_I_9.csv (Signal Sufficiency) — the pattern_name field only	The only narrow case worth embedding on this sheet — a rough early guess at which of ~185 patterns applies, before the exact signal checklist takes over	Embedded; used for a first coarse match, then the exact minimum/must-have signal check (Bucket 3) takes over
Kb_Phase_I_7.csv (Red Flags) — the trigger_phrases column	Real examples, deliberately written to include disguised/indirect language — a genuine "does this match, by meaning" problem	Embedded with a deliberately more permissive match threshold than normal (see 08_SAFETY_INTERLOCK_AND_TRIAGE.md §4) — but see the special case below for what happens after a match
PHASE II Emotion Associated user expression.docx , List of Emotions - Tanisha.docx , and the Pathway... doc's emotion picklist — once merged into one canonical list (see §Blockers below)	Real first-person "I feel..." phrasing, meant to teach/pattern-match emotion detection	Embedded and matched against a user's message to tag which emotion is present
List of Scenarios - Tanisha.docx , Anxiety Situations.docx	Master scenario/topic taxonomies — matching "which life situation is this" is a real semantic-similarity problem, not a lookup	Embedded; used to route a conversation toward the right topic-specific question bank or conversation pattern
Layer I Questions and Statments.docx , Layer II Question Bank.docx , Layer III.docx , PHASE II Layer I and IV Questions.docx , Layer 2 to 3 Transitions.docx	These are a menu to select the single best-fitting phrase from, given the current conversational depth — a retrieval problem, not free generation	Embedded per layer; the pipeline picks the closest-fitting question/statement for "what layer are we in right now" and hands it to the LLM as phrasing material (never the LLM inventing the question from scratch)

2b. Recommending (real content handed back to the user)

File (exact name)	Why it belongs here	How it's actually used
dbt-skills-workbook.docx / dbt-skills-workbook-simple.pdf	A finished, ready-to-use DBT skills tool — literally the "self-care tool" content 04_CONVERSATION_PIPELINE.md Step F was designed to search for and previously had nothing to find	Chunked (one skill = one chunk), embedded, and genuinely semantically searched at Tier 1 — the one part of the whole system where full RAG search is used exactly as designed
brain-dump-worksheets.pdf	Same reasoning — finished, ready-to-hand-out content	Same mechanism as the DBT workbook

File (exact name)	Why it belongs here	How it's actually used
SELF-HELP CONVERSATIONS.docx (the toolkit half — Behavioral Activation, Time Management, Mindfulness/Grounding, Journaling, Tracking Worksheets sections)	Broadest single self-help asset received — real, usable exercises	Same mechanism; each named tool/worksheet becomes its own searchable chunk
Phase III/IV content in general, once the client's remaining gaps are filled (see 15_OPEN_QUESTIONS_AND_BLOCKERS.md #24)	This is the category Phase III (clinical guidance/psycho-education) and Phase IV (self-help tools) were always meant to populate	Same mechanism — this bucket grows as more finished content arrives

Bucket 3 — Becomes literal code logic (never text, never embedded, never searched)

File (exact name)	Why it belongs here	How it's actually used
Kb_Phase_I_C11.csv (Routing Rules)	A fixed, precise signal-combination → tier mapping — searching for it would be slower and less reliable than reading it off a table	Loaded as a direct lookup table, evaluated strictly in priority column order (1→2→3), never rule_id order (see 05_KNOWLEDGE_BASE_DEEP_DIVE.md §4) — app/safety/routing.py in the backend
Kb_Phase_I_9.csv (Signal Sufficiency) — every field except pattern_name (minimum_signals , must_have_signals , confidence_rule , graceful_exit_trigger)	Exact lookup criteria, checked once a candidate pattern is already known from Bucket 2a's coarse match	Read directly off the row matched in 2a — no second search
Kb_Phase_I_17.csv (Fallback & Safety Net) — the trigger_condition field	Describes the conversation's own shape (turn count, response length, confidence score) — not something a user says, so there's no phrase to search for	Evaluated as code against the running conversation state (Redis) every turn — app/safety/fallback.py

File (exact name)	Why it belongs here	How it's actually used
Pathway For AI In Different Situations - Tanisha.docx	This is, almost verbatim, a 15-step decision flow with explicit branches — the closest thing to pseudocode the client has produced	Becomes the actual control flow of app/pipeline/orchestrator.py — identify concern → categorize → safety check → assess emotion/severity → pick intervention → close. Its vocabulary (Severity buckets, the 22-item category list) still needs reconciling with the existing Tier/scenario systems first — see 15_OPEN_QUESTIONS_AND_BLOCKERS.md #23
Knowledge_Base_Directory	Internal build-planning document only	Never reaches the AI at all — not ingested in any form
Filler-phrase library (Mindfully_Yours_AI_Latency_Filler_Phrases.pdf)	Selection is a rule match (situation type + expected gap length → category), not a meaning-based search	A lookup table in app/pipeline/filler.py , picked by rule, never embedded

Special case: exact-fetch content (a hybrid of Bucket 2's matching and Bucket 3's "no thinking allowed")

Some content is found via a Bucket 2a match, but the response itself must be used **word-for-word**, never paraphrased by the LLM — this is its own category precisely because getting it wrong is the single biggest safety risk in the system:

File (exact name)	Field	Rule
Kb_Phase_I_7.csv (Red Flags)	bot_script	Matched via embedding (Bucket 2a), then fetched verbatim — the LLM is never called for this response. See 08_SAFETY_INTERLOCK_AND_TRIAGE.md §4 for the 4 flags (RF-028-031) that currently have no script at all and must hard-alert a human instead.
Kb_Phase_I_17.csv (Fallback & Safety Net)	bot_script	Same rule — fetched verbatim once a fallback scenario is confirmed by Bucket 3's code check.

What's explicitly NOT usable yet (don't wire these in)

File (exact name)	Status	Why
Career Conversation Sets 1.docx	Not real variety	1 conversation template mechanically repeated ~500 times with only a career-field name swapped — the underlying list of ~500 career names is the only reusable part
Anam/Conversation Combinations-sleep.docx , Anam/Conversations Combinations- anxiety 3rd person.docx , Anam/Miscellaneous Conversations.docx	Still password-locked	Didn't decrypt with the password used for the rest of the drop — need a second password from Anam
Paraphilic-disorder content in List of Scenarios - Tanisha.docx ("Fetishes" section) and the Week 2 conversation-script file	On hold	Needs an explicit clinical + legal governance decision before it goes anywhere near a live index — see 15_OPEN_QUESTIONS_AND_BLOCKERS.md #26
Conversations - Depression.docx , Conversations Combination (Bipolar and Psychosis).docx , Shreya/Trauma Related Conversations.docx , Conversation Combinations- Life Transitions.docx , Ankita/Conversations around Parenting.docx	Partially usable	Each has real, specific content gaps (missing scenarios, cuts off mid-file) — usable for what's actually there, but don't treat as complete; see 15_OPEN_QUESTIONS_AND_BLOCKERS.md #24

Blockers this data-usage map depends on being resolved first

Before Bucket 2a's emotion-matching and Bucket 3's Pathway-as-code work can be built for real (not just structurally stubbed), these need a decision — full detail in

[15_OPEN_QUESTIONS_AND_BLOCKERS.md](#):

- **#23** — merge the three unreconciled emotion lists (19/9/11 labels) into one canonical taxonomy, and reconcile the new Severity (Mild/Moderate/Severe/Crisis) scale and 22-item category list against the existing Tier 1/2/3 and ~9-category systems.
- **#25** — standardize tier-labeling, since different clinical-team contributors used it inconsistently.

Companion to [04_CONVERSATION_PIPELINE.md](#), [05_KNOWLEDGE_BASE_DEEP_DIVE.md](#), and [18_NEW_CONVO_DATA_CATALOG.md](#). Update this table the moment a new file arrives or a blocker above is resolved — this is meant to be the one place a developer checks before deciding how to treat a new piece of content, not a snapshot to go stale.